

Text-To-SQL App with feedback loop

Natural Language to SQL Query Generation using LangChain & Ollama

BDNSV - Dima Florin-Alexandru (462)

FMI Unibuc - 2026

Project Overview

What is project's purpose?

- ▶ **Goal:** Convert natural language questions into SQL queries
- ▶ **Multi-step chain:**
 - ▶ Accept NLP question
 - ▶ Generate SQL query using feedback loop for learning from user interactions
 - ▶ Execute query
 - ▶ Summarize results in natural language
- ▶ **Tech Stack:** LangChain, Ollama (Local LLMs), SQLite, Gradio (and more)



Core Features

Meeting project requirements

- ▶ **1. Ambiguous Input Handling:**
 - ▶ Schema-aware prompt engineering
 - ▶ Database context inclusion
 - ▶ Clarification requests for vague questions
- ▶ **2. LLM Comparison:**
 - ▶ Llama3, Mistral, CodeLlama support
 - ▶ Side-by-side accuracy evaluation
 - ▶ Syntax and logic comparison
- ▶ **3. Error Handling & Feedback:**
 - ▶ SQL validation and auto-correction
 - ▶ User feedback loop (1-5 rating + corrections)
 - ▶ Few-shot and chain-of-thought prompting
- ▶ **4. Benchmark Testing:**
 - ▶ Spider benchmark: measure exact match, execution accuracy


System Architecture

Multi-step pipeline & Data flow



System Architecture

Correction & Learning activation



Feedback loop for self-improving

How does it work?

- ▶ **Step 1: Collect Feedback**
 - ▶ User rates query (1-5 stars)
 - ▶ Optional: provide corrected SQL
- ▶ **Step 2: Learn**
 - ▶ Extract positive examples (rating ≥ 4)
 - ▶ Store user corrections
 - ▶ Find similar past queries
- ▶ **Step 3: Improve**
 - ▶ Enhance prompts with learned examples
 - ▶ Add correction guidance
 - ▶ Better results → More positive feedback
- ▶ **Result:** System improves automatically with usage

LLM Comparison

Multi-Model evaluation

- ▶ **Models Tested:** Llama3:latest, Mistral:7b, CodeLlama:latest
- ▶ **Metrics:**
 - ▶ Query accuracy
 - ▶ Syntax correctness
 - ▶ Logical structure
 - ▶ Execution success rate
 - ▶ Response time
- ▶ **Finding:**
 - ▶ CodeLlama performs best for complex queries
 - ▶ Llama3 performs best for natural language understanding

Benchmark results

Spider benchmark performance

- ▶ **Dataset:** Spider
- ▶ **Initial results (no learning):** Exact match: 20%, Execution accuracy: 35%
- ▶ **After normalization & Learning:** Exact match: 75%, Execution accuracy: 75%
- ▶ **Key Improvements:**
 - ▶ SQL normalization (aliases, whitespace, case)
 - ▶ Few-shot learning from positive examples
 - ▶ Error correction from user feedback
- ▶ **Test set:** 5-20 random queries across multiple databases

User Interface

Dual interface: CLI & Web

- ▶ **1. CLI Interface:**
 - ▶ Interactive command-line
 - ▶ Commands: stats, learning, compare, quit
 - ▶ Real-time feedback collection
 - ▶ Correction submission for failed queries
- ▶ **2. Web Interface (Gradio):**
 - ▶ User-friendly visual interface
 - ▶ Tabs: Query Generation, Model Comparison, Learning System
 - ▶ Real-time feedback collection
 - ▶ Correction submission for failed queries
- ▶ **Accessibility:** Both interfaces suitable for different user preferences

Conclusion

Summary & Next steps

- ▶ **Achievements:**
 - ▶ Text-to-SQL pipeline meeting the requirements
 - ▶ Feedback loop with automatic learning
 - ▶ Fair results achieved on Spider benchmark
 - ▶ Multi-LLM comparison framework
 - ▶ Dual interface (CLI + Web)
 - ▶ Error handling
- ▶ **Documentation:** Comprehensive README, guides, diagrams, code comments
- ▶ **Possible future enhancements:**
 - ▶ Support for more complex SQL (subqueries, CTEs)
 - ▶ Multi-database support (PostgreSQL, MySQL)
 - ▶ Advanced learning algorithms/fine-tuning
 - ▶ Query optimization suggestions
 - ▶ More benchmarks (BIRD, WikiSQL)

References

References & Resources

- ▶ **AI Tools Used:**
[1] GitHub Copilot (Claude Sonnet 4.5), <https://github.com/features/copilot>, Date generated: January 2026
- ▶ **Frameworks & Libraries:**
[2] LangChain Documentation, <https://python.langchain.com/>, Last accessed: January 2026
[3] Ollama, <https://ollama.ai>, Last accessed: January 2026
- ▶ **Benchmarks:**
[4] Yu, T., et al., "Spider: A Large-Scale Human-Labeled Dataset for Complex and Cross-Domain Semantic Parsing and Text-to-SQL Task", EMNLP 2018
- ▶ **Database:**
[5] SQLite, <https://www.sqlite.org>, Last accessed: January 2026
- ▶ **Code Repository:**
[6] GitHub Repository, <https://github.com/dimaflorinalex/ProjectBDNSV>